

DYNAMIC ENSEMBLE LEARNING FOR CREDIT CARD FRAUD DETECTION - CRISP-DM PROCESS REPORTs

BUSINESS UNDERSTANDING

DATA UNDERSTANDING: *Data Description Report*

- **Location:** The original dataset is available at <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>
- **Method used to acquire:** Download the dataset using the kagglehub library, which provides a simple way to interact with Kaggle resources.
- **Format:** csv file
- **Quantity:** Rows (284807) - Columns (31)
- **Identities of the fields:** The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two day. It contains only numerical input variables which are the result of a PCA transformation. Due to confidentiality issues, we cannot provide the original features and more background information. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset.
- **Surfaces features:** Amount, Time, PCs (V1 ... V28), Class
 - **Quantitative Variables**
 - Discrete: -
 - Continuous: Amount, Time, PCs (V1 ... V28)
 - **Categorical Variables**
 - Nominal: Class (0=Legit / 1=Fraud)
 - Ordinal: -

DATA UNDERSTANDING: *Data Quality Report*

- **Missing values Analysis**
 - **Column-Level Analysis:** Checked the presence of missing values in each of the columns to understand if some imputation was required. There are 0 columns with missing values.
 - **Row-Level Analysis:** Checked the presence of rows with more than 50% of missing values among the columns. There are 0 rows with more than 50% of missing values among the columns.
- **Class Balance:** The dataset presents a very high imbalance, as already know for frauds classification problems: Legit-0 (284315 – 99.827%) / Fraud (492 – 0.173%). In the next steps the data imbalance should be handled by using undersampling or oversampling techniques or by adopting models that can handle this very high imbalance within the dataset.
- **Duplicated Row:** The dataset contains 1081/284807 (0.380%) duplicated rows. In particular, it contains 773 duplicated patterns distributed as follows (611 different patterns repeated 2 times, 66 different patterns repeated 3 times, 81 different patterns repeated 4 times, 10 different patterns repeated 5 times, 1 pattern repeated 6 times, 2 different patterns repeated 9 times, 2 different patterns repeated 18 times). Additionally, we check duplicated both in terms of all features + Class and only looking to the features; in both case we have 1081/284807 (0.380%).

If we leave duplicates in the dataset we may end up in:

 - Bias in the model - The classifier will “see” repeated evidence of the same row. This gives unfair influence to those examples during training.
 - Risk of overfitting - Especially with a small number of fraud cases, repeating exact Class 1 samples inflates the model's confidence.
 - Invalidation of cross-validation - Duplicates might leak into both training and validation folds. This leads to optimistically biased performance metrics (especially precision/recall).

Based on these considerations, one solution could be of dropping duplicated rows, keeping each of them just one time.

DATA UNDERSTANDING: *Data Exploration Report*

- **Class Balance:** The dataset presents a very high imbalance: Legit-0 (284315 – 99.827%) / Fraud (492 – 0.173%).
- **Time Analysis:** Time contains the seconds elapsed between each transaction and the first transaction in the dataset
 - The skewness value of approximately -0.036 is very close to zero. This indicates that the distribution is highly symmetrical.
 - The mean (94813.86) is slightly greater than the median (84692.00). This small difference confirms the slight negative skew (mean should be less than the median for negative skew) and indicates that the data has a slightly heavier tail on the positive (right) side than on the negative (left) side, despite the skewness value being negative.
 - The time values range from 0.00 to 172792.00. The middle 50% of the data spans from 54201.50 (25%) to 139320.50 (75%).
 - The standard deviation (47488.14) is relatively large compared to the mean, but the symmetrical nature suggests the spread is consistent with a normal-like distribution, and the min/max values appear reasonable for the overall range.
 - For Legit transactions, we can observe a repetitive wave pattern with two large peaks. This suggests transaction volume follows a diurnal (daily) cycle: possibly office hours and nighttime dips, repeated over two days. This gives us confidence that time reflects real-world temporal behavior.
 - For Fraud transactions, we have much lower volume, but the key insight is that fraud attempts are not uniformly distributed. We can see clusters/spikes in specific time windows. This suggests frauds may be targeted in bursts, likely by bots or coordinated attacks.
- **Transaction Amount Distribution**
 - The distribution of Amount is highly skewed right (approximately 16.98). This means the vast majority of the transactions have very small amounts, and the long, thin tail of the distribution extends far out to the right due to a small number of extremely large transactions. These large transactions are acting as significant outliers.
 - The mean (88.47) is far greater than the median (22.00). The mean is being pulled significantly upwards by the extremely large values in the right tail, while the median remains a robust measure of the typical transaction amount.
 - The standard deviation (250.4) is extremely large compared to the mean. This indicates a long right tail, which is a common indicator of a highly skewed distribution and the presence of extreme outliers.
 - Most transactions are small (75% of transactions are ≤ 77.51).
 - There are zero-amount transactions. This may represent failed/incomplete attempts, errors, or may indicate fraud tests (e.g., verifying a card).
 - Logarithmic Transformation [$\log(x + 1)$] is applied, end up into a more symmetric and usable form. This transformation has successfully normalized the Amount feature, pulling the extreme right tail back towards the center of the distribution. Log transform stabilizes variance, compresses outliers, and makes fraud/legit distributions visually and statistically manageable.
- **PCA Feature Analysis**
 - We don't have information on how the PCA was computed. Thus, we perform preliminary checks knowing how the PCA works and what is expected:
 - PCA must be performed on centered data (mean=0).
 - All the PCs (V1 ... V28) have mean ~ 0 , suggesting that data have been centered before applying the PCA (also the kde plots show that the distribution are centered around 0);
 - In PCA we have a variance reduction from the first to the last principal direction.
 - We compute the standard deviation and we can confirm that there is a variance reduction, as expected.
 - Median and quartile are symmetric distributed and centered.

- Correlation matrix shows that all the PCs have 0 correlation, as expected, since PCA applies orthogonal projections.
 - We use the cumulative explained variance to find a cut-off and find the number of PCs to retain at least 85% (V1 ... V17), 90% (V1 ... V20) and 95% (V1 ... V23) of the total variance.
 - Comparing the kde for each PCs between the two classes emerged:
 - V4, V10, V11, V12, V14, V17 are highly promising candidates for model discrimination and maybe feature selection.
 - V2, V3, V5, V6, V8, V13, V16, V18, V22 may not separate well alone but could help in interaction with others.
 - V1, V7, V9, V19, V20, V21, V23, V28, V15, V26 are probably noise or redundant for classification.
- **Features Relevance to Target class**
 - ***Pearson Correlation***
 - PCA features V11–V17 carry strong linear separability.
 - V11, V12, V14, V17, V10 showed strongest (negative) correlation with fraud class.
 - LogAmount, Hour, Time showed weak or no linear correlation.
 - Amount close to 0 correlation.
 - ***Mutual Information***
 - No feature has extremely high MI (> 0.01), which is expected due to high class imbalance and feature obfuscation (PCA).
 - V17, V14, V12, V11, V10, V27 showed highest MI — top fraud indicators.
 - LogAmount, Hour — Moderate MI (more relevant than many PCA features).
 - Amount (raw) — Less informative than LogAmount.
 - Many PCA features (e.g., V1, V4, V3, V8) — Carry small but non-zero signal.

DATA PREPARATION: *Data Sampling Report*

- **Sampling Approach**
 - ***keep_all_minority_with_ratio***: To address the class imbalance in the 284315 record dataset [Legit-0 (284315 – 99.827%) / Fraud (492 – 0.173%)], a sub-sampling strategy was employed to create a smaller, computationally efficient, and representative sample. The strategy, "*keep_all_minority_with_ratio*", involved retaining all instances of the minority 'Fraud' class and performing a proportional random under-sampling of the 'Legit' class at a 50:1 ratio. This process yielded a new sample of 25092 records [Legit-0 (24600 – 98.039%) / Fraud (492 – 1.961%)].
 - Running configuration:
 - *policy='keep_all_minority'*
 - *seed=42*
 - *n_rows=None*
 - *frac=None*
 - *ratio=50*
- **Post Sampling Validation**
 - ***A rigorous two-part statistical validation*** was then conducted to ensure this new sample remained representative of the original population. The validation confirms that the sub-sampled dataset is a statistically robust and highly representative proxy for the original data.
 - ***Validation of Individual Feature Distributions (Univariate Analysis)***:
 - We first tested whether the distribution of each individual feature in the sample (n=25092) was statistically identical to its distribution in the population (N=284315).
 - Initial checks using descriptive statistics (mean, median, quartiles, std. dev) and visualizations (histograms, violin plots) showed a near-perfect visual overlap for all features.
 - A two-sample Kolmogorov-Smirnov (K-S) test was formally conducted for all features.
 - ***Validation of Feature Relationships (Multivariate Analysis)***:
 - We then tested whether the complex relationships between features were preserved during sampling.
 - We computed the Pearson correlation matrix for the original population and the new sample.
 - An "error matrix" was calculated by taking the absolute difference between the two correlation matrices.
 - The analysis of this error matrix was excellent.
 - Mean Error: The average error across all feature relationships was approximately 0.0105 (1.05%).
 - Max Error: The single "worst-case" error observed for any feature pair was approximately 0.102 (10.2%).

The K-S tests confirm that individual features retained their statistical properties. The correlation analysis confirms that the relationships between features were also preserved with extremely high fidelity (99% average accuracy).

DATA PREPARATION: *Data Cleaning Report*

In the Data Quality phase we assess the following quality problems:

- **Missing values Analysis**
 - ***Column-Level Analysis***: no-problems
 - ***Row-Level Analysis***: no-problems
- **Class Balance**: The initial dataset presents a very high imbalance, as already know for frauds classification problems: Legit-0 (284315 – 99.827%) / Fraud (492 – 0.173%). After the data sampling phase (performed to reduce the dataset size for computational limits), the new dataset still have imbalance [Legit-0 (24600 – 98.039%) / Fraud (492 – 1.961%)].
To avoid leakage and bias, we choose to use oversampling techniques to rebalance the dataset, which will be implemented during the training phase only on the training dataset (and not on the validation or test datasets).

- **Duplicated Row:** The dataset contains 1081/284807 (0.380%) duplicated rows in the initial dataset. To avoid bias in the model – since the classifier will “see” repeated evidence of the same row – we choose to remove these duplicated occurrences.
The sampled dataset contains 25092 rows [Legit-0 (24600 – 98.039%) / Fraud (492 – 1.961%)]. We found 25/25092 (0.10%) of duplicated rows. After dropping the cleaned dataset presents the following statistics: 25067 rows [Legit-0 (24594 – 98.113%) / Fraud (492 – 1.887%)]

DATA PREPARATION: *Data Construction Report*

- **Amount:** In the Data Exploration phase we assess that the Amount feature is highly right skewed (with very long tail). Moreover, there are zero-amount transactions. This may represent failed/incomplete attempts, errors, or may indicate fraud tests (e.g., verifying a card).
 - **Logarithmic Transformation** [$\log(x + 1)$] is applied, end up into a more symmetric and usable form. This transformation has successfully normalized the Amount feature, pulling the extreme right tail back towards the center of the distribution. Log transform stabilizes variance, compresses outliers, and makes fraud/legit distributions visually and statistically manageable. This transformation can be directly applied with no-leakage risk.
 - **StandardScaler** is then applied on Amount_log1p to transform the distribution into a Standard distribution with 0 mean and unit variance. This step is not applied on the dataset but it is part of the training pipeline. This transformation is always fit only on the training dataset and then the computed statistics are used to transform validation and test data. In this way we can avoid data leakage from the test set.
- **Time:** Time contains the seconds elapsed between each transaction and the first transaction in the dataset. We chose to transform this feature to better represent cyclicity.
 - **Trigonometric features** are used to encode Time periodic feature using a sine and cosine transformation with the matching period. Each ordinal time feature is transformed into 2 features that together encode equivalent information in a non-monotonic way, and more importantly without any jump between the first and the last value of the periodic range.
 - **Period** Since we have two days of data in seconds, we have $60 \times 60 \times 24 \text{ sec} = 86400 \text{ sec}$ in one day. We use period = 86400 to represent these two days of data in terms of sine and cosine on the trigonometric circle.
- **PCA Feature V1 ... V28:** As already checked in previous steps, these are the PCs already computed (with all the required adjustment computed). No additional transformation are required.